

HUMAN VIDEO PHYSICAL INTENT DISAMBIGUATION: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study learning robots from human video. The motivating bet is: Separate human intent from morphology-specific motion artifacts in video demonstrations.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For learning robots from human video, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2003; 2025a; 2018; 2022; 2009; 2025b). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.515	0.485
observed_only_wam	hidden_mode_shift	0.177	0.823
augmentation_as_data_wam	nominal	0.615	0.385
augmentation_as_data_wam	hidden_mode_shift	0.245	0.755
scalar_uncertainty_wam	nominal	0.642	0.358
scalar_uncertainty_wam	hidden_mode_shift	0.325	0.675
proposed_branch_mechanism	nominal	0.772	0.228
proposed_branch_mechanism	hidden_mode_shift	0.578	0.422

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Human Video Physical Intent Disambiguation argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Estimating articulated human motion with covariance scaled sampling, 2003. <https://doi.org/10.1177/0278364903022006003>.
- A probabilistic model of human motion and navigation intent for mobile robot path planning, 2009. <https://doi.org/10.1109/icara.2000.4803931>.
- Human motion intent learning based motion assistance control for a wearable exoskeleton, 2018. <https://doi.org/10.1016/j.rcim.2017.08.007>.
- A bio-inspired mechanism for learning robot motion from mirrored human demonstrations, 2022. <https://doi.org/10.3389/fnbot.2022.826410>.
- Cori: Communication of robot intent for physical human-robot interaction, 2025a. <http://arxiv.org/abs/2505.20537v2>.
- Replay: Robot embodiment via intent-aware policy imitation by replicating human demonstrations from video, 2025b. <https://doi.org/10.1007/s12555-025-0505-8>.